

# Named Entity Recognition

Author: Christin Seifert, licensed under the Creative Commons Attribution 3.0 Unported License <https://creativecommons.org/licenses/by/3.0/>

This is a tutorial for NER (named entity recognition). In this tutorial you will see

- how to apply a pre-trained named entity recognition model to your text

It is assumed that you have some general knowledge on

- .. no particular knowledge required. You should be able to read texts, though ;-)

**Prerequisites.** We first need to install the Stanford NER tagger from [here](#). And java also has to be installed. You have to figure out

- where the jar file `stanford-ner.jar` is located
- where the pretrained models (e.g. `english.all.3class.distsim.crf.ser.gz`) is located, this is the subdirectory classifiers
- whether the right version of java is installed. On a command line type `java -version` to see the version. Refer back to the documentation on the stanford nlp page to see which version is needed.

You can also test the NER tagger online [here](#).

In [7]:

```
from nltk.tag import StanfordNERTagger
from nltk.tokenize import word_tokenize

# Adapt those lines to your installation
jar_location = '/Users/sech/stanford-ner-2018-10-16/stanford-ner.jar
model_location_3classes = '/Users/sech/stanford-ner-2018-10-16/classi
model_location_7classes = '/Users/sech/stanford-ner-2018-10-16/classi
st3 = StanfordNERTagger(model_location_3classes, jar_location, encoding
```

```
st7 = StanfordNERTagger(model_location_7classes, jar_location, encoding='utf-8')

print(st3)
print(st7)
```

```
<nltk.tag.stanford.StanfordNERTagger object at 0x1a1d98af60>
<nltk.tag.stanford.StanfordNERTagger object at 0x1a1d98af28>
```

Let's take a paragraph from the [Wikipedia page of Ada Lovelace](#) as an example. We need to put the text in triple quotes since the text itself contains quoting characters.

In [10]:

```
text = '''Lovelace became close friends with her tutor Mary Somerville'''
print(text)
```

```
Lovelace became close friends with her tutor Mary Somerville, who int
```

First we need to tokenize the text and then we apply the NER tagger. Let's try both, the 3 class version and the 7 class version.

In [11]:

```
tokenized_text = word_tokenize(text)
text_ner3 = st3.tag(tokenized_text)
text_ner7 = st7.tag(tokenized_text)

print(text_ner3)
print(text_ner7)
```

```
[('Lovelace', 'PERSON'), ('became', '0'), ('close', '0'), ('friends',
[('Lovelace', '0'), ('became', '0'), ('close', '0'), ('friends', '0')]
```

We see that each word is tagged. Tags are for instance `ORGANIZATION` or `PERSON`. Very prominently, the `0` tag appears often. This is the `other` class (everything that is not an organisation or person, etc.). But it is still an awful lot of text. Let's just have a look at the non-other entities detected. We do this assuming that adjacent words having the same tag should be collapsed into one named entity.

In [13]:

```
from itertools import groupby

print("**** 3 classes ****")
for tag, chunk in groupby(text_ner3, lambda x:x[1]):
    if tag != "0":
        print("%-12s"%tag, " ".join(w for w, t in chunk))

print("**** 7 classes ****")
for tag, chunk in groupby(text_ner7, lambda x:x[1]):
    if tag != "0":
        print("%-12s"%tag, " ".join(w for w, t in chunk))
```

```
**** 3 classes ****
PERSON      Lovelace
PERSON      Mary Somerville
PERSON      Charles Babbage
LOCATION      Somerville
PERSON      Andrew Crosse
PERSON      David Brewster
PERSON      Charles Wheatstone
PERSON      Michael Faraday
PERSON      Charles Dickens
PERSON      Ada
PERSON      John Hobhouse
PERSON      Byron
PERSON      Ada
```

```
PERSON      Hobhouse
**** 7 classes ****
PERSON      Mary Somerville
PERSON      Charles Babbage
DATE        1833
LOCATION      Somerville
PERSON      Andrew Crosse
PERSON      David Brewster
PERSON      Charles Wheatstone
PERSON      Michael Faraday
PERSON      Charles Dickens
DATE        1834
ORGANIZATION Ada
PERSON      John Hobhouse
PERSON      Byron
DATE        February 1834
ORGANIZATION Ada
PERSON      Hobhouse
```

We see that while this is pretty impressive, it still makes errors. For example, one occurrence of Ada is tagged a `ORGANISATION`. You should take the non-perfect nature into account if you use those tags further in your nlp pipeline.

That's all.

In [ ]:

---

Content source: [chseifert/tutorials](https://notebook.community/chseifert/tutorials/nlp-ie/Named-Entity-Recognition)

Similar notebooks:

- [Named-Entity-Recognition](#)
- [class14](#)

- [Text-Preprocessing](#)
  - [HW7](#)
  - [fizzbuzz two ways](#)
  - [inherit](#)
  - [test\\_tensorflow\\_and\\_jupyter\\_install](#)
  - [UsingNERCs](#)
  - [keras\\_preprocessing\\_image](#)
  - [Sunoikisis - Named Entity Extraction](#)
- [1b-VV](#)

[notebook.community](#) | [gallery](#) | [about](#)